

```
from tensorflow import keras
from keras.datasets import imdb
%matplotlib inline
import numpy as np
import pandas as pd
from matplotlib import cm
import matplotlib.pyplot as plt
import seaborn as sns
import os
import time

from keras.preprocessing import sequence
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation
from keras.layers import Embedding
from keras.layers import Conv1D, GlobalMaxPooling1D
from keras.callbacks import EarlyStopping
from keras import models

(X_train, y_train), (X_test, y_test) = imdb.load_data()
X = np.concatenate((X_train, X_test), axis=0)
y = np.concatenate((y_train, y_test), axis=0)

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 [=====] - 0s 0us/step

##training data shape review
print("Training data: ")
print(X.shape)
print(y.shape)
print("Classes: ")
print(np.unique(y))

Training data:
(50000,)
(50000,)
Classes:
[0 1]

print("Number of words: ")
print(len(np.unique(np.hstack(X))))

Number of words:
88585

print("Review length: ")
result = [len(x) for x in X]
print("Mean %.2f words (%f)" % (np.mean(result), np.std(result)))
# plot review length
plt.boxplot(result)
plt.show()
```

Review length:
Mean 234.76 words (172.911495)

```
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=5000)

def vectorize_sequences(sequences, dimension=5000):
    # Create an all-zero matrix of shape (len(sequences), dimension)
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1. # set specific indices of results[i] to 1s
    return results

# Our vectorized training data
x_train = vectorize_sequences(train_data)
# Our vectorized test data
x_test = vectorize_sequences(test_data)

# Our vectorized labels one-hot encoder
y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')

from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Dense(32, activation='relu', input_shape=(5000,)))
model.add(layers.Dense(32, activation='relu',))
model.add(layers.Dense(1, activation='sigmoid'))

#Set validation set aside

x_val = x_train[:10000]
partial_x_train = x_train[10000:]

y_val = y_train[:10000]
partial_y_train = y_train[10000:]

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['acc'])

start_time_m1 = time.time()
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))
total_time_m1 = time.time() - start_time_m1

Epoch 1/20
30/30 [=====] - 3s 47ms/step - loss: 0.5519 - acc: 0.7599 - val_loss: 0.3860 - val_acc: 0.8554
Epoch 2/20
30/30 [=====] - 2s 52ms/step - loss: 0.2953 - acc: 0.8919 - val_loss: 0.3018 - val_acc: 0.8791
Epoch 3/20
30/30 [=====] - 1s 46ms/step - loss: 0.2138 - acc: 0.9210 - val_loss: 0.2864 - val_acc: 0.8847
Epoch 4/20
30/30 [=====] - 1s 37ms/step - loss: 0.1697 - acc: 0.9391 - val_loss: 0.3020 - val_acc: 0.8773
Epoch 5/20
30/30 [=====] - 1s 30ms/step - loss: 0.1386 - acc: 0.9517 - val_loss: 0.3249 - val_acc: 0.8766
Epoch 6/20
30/30 [=====] - 1s 32ms/step - loss: 0.1165 - acc: 0.9607 - val_loss: 0.3521 - val_acc: 0.8724
Epoch 7/20
30/30 [=====] - 1s 33ms/step - loss: 0.0943 - acc: 0.9703 - val_loss: 0.3855 - val_acc: 0.8717
Epoch 8/20
30/30 [=====] - 1s 32ms/step - loss: 0.0722 - acc: 0.9799 - val_loss: 0.4203 - val_acc: 0.8635
Epoch 9/20
30/30 [=====] - 1s 33ms/step - loss: 0.0542 - acc: 0.9887 - val_loss: 0.4556 - val_acc: 0.8613
Epoch 10/20
30/30 [=====] - 1s 32ms/step - loss: 0.0413 - acc: 0.9931 - val_loss: 0.4969 - val_acc: 0.8611
Epoch 11/20
30/30 [=====] - 1s 33ms/step - loss: 0.0302 - acc: 0.9961 - val_loss: 0.5281 - val_acc: 0.8616
Epoch 12/20
30/30 [=====] - 1s 31ms/step - loss: 0.0205 - acc: 0.9983 - val_loss: 0.5751 - val_acc: 0.8594
Epoch 13/20
```

```

30/30 [=====] - 1s 34ms/step - loss: 0.0141 - acc: 0.9997 - val_loss: 0.6076 - val_acc: 0.8600
Epoch 14/20
30/30 [=====] - 1s 35ms/step - loss: 0.0095 - acc: 0.9999 - val_loss: 0.6441 - val_acc: 0.8601
Epoch 15/20
30/30 [=====] - 1s 47ms/step - loss: 0.0066 - acc: 1.0000 - val_loss: 0.6822 - val_acc: 0.8599
Epoch 16/20
30/30 [=====] - 2s 58ms/step - loss: 0.0048 - acc: 1.0000 - val_loss: 0.7112 - val_acc: 0.8597
Epoch 17/20
30/30 [=====] - 1s 28ms/step - loss: 0.0036 - acc: 1.0000 - val_loss: 0.7393 - val_acc: 0.8596
Epoch 18/20
30/30 [=====] - 1s 31ms/step - loss: 0.0028 - acc: 1.0000 - val_loss: 0.7646 - val_acc: 0.8598
Epoch 19/20
30/30 [=====] - 1s 35ms/step - loss: 0.0022 - acc: 1.0000 - val_loss: 0.7878 - val_acc: 0.8594
Epoch 20/20
30/30 [=====] - 1s 27ms/step - loss: 0.0018 - acc: 1.0000 - val_loss: 0.8095 - val_acc: 0.8592

```

```
print("The Dense Convolutional Neural Network 1 layer took %.4f seconds to train." % (total_time_m1))
```

The Dense Convolutional Neural Network 1 layer took 43.5509 seconds to train.

```

history_dict = history.history
history_dict.keys()

```

```
dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])
```

```

import matplotlib.pyplot as plt
%matplotlib inline

```

```

acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

```


```
epochs = range(1, len(acc) + 1)
```

```

# "bo" is for "blue dot"
plt.plot(epochs, loss, 'bo', label='Training loss')
# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

```
plt.show()
```



```


plt.clf() # clear figure
acc_values = history_dict['acc']
val_acc_values = history_dict['val_acc']

plt.plot(epochs, acc, 'bo', label='Training acc')

```

```
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



```
model.summary()
```

```
Model: "sequential"
Layer (type)                Output Shape              Param #
=====
dense (Dense)                (None, 32)                160032
dense_1 (Dense)              (None, 32)                1056
dense_2 (Dense)              (None, 1)                 33
=====
Total params: 161,121
Trainable params: 161,121
Non-trainable params: 0
```

```
from sklearn.metrics import confusion_matrix, accuracy_score, auc
#predictions
pred = model.predict(x_test)
classes_x=np.argmax(pred,axis=1)

#accuracy
accuracy_score(y_test,classes_x)

782/782 [=====] - 3s 2ms/step
0.5

#Confusion Matrix
conf_mat = confusion_matrix(y_test, classes_x)
print(conf_mat)

conf_mat_normalized = conf_mat.astype('float') / conf_mat.sum(axis=1)[:, np.newaxis]
sns.heatmap(conf_mat_normalized)
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
[[12500    0]
 [12500    0]]
Text(0.5, 23.52222222222222, 'Predicted label')
```



```
#Dense with Two Layer
model2 = models.Sequential()
model2.add(layers.Dense(32, activation='relu', input_shape=(5000,)))
model2.add(layers.Dense(32, activation='relu'))
model2.add(layers.Dense(32, activation='relu'))
model2.add(layers.Dense(1, activation='sigmoid'))

model2.compile(optimizer='adam',
               loss='binary_crossentropy',
               metrics=['acc'])

start_time_m2 = time.time()
history= model2.fit(partial_x_train,
                   partial_y_train,
                   epochs=20,
                   batch_size=512,
                   validation_data=(x_val, y_val))
total_time_m2 = time.time() - start_time_m2

print("The Dense Convolutional Neural Network 2 layers took %.4f seconds to train." % (total_time_m2))

Epoch 1/20
30/30 [=====] - 2s 39ms/step - loss: 0.5316 - acc: 0.7702 - val_loss: 0.3490 - val_acc: 0.8643
Epoch 2/20
30/30 [=====] - 1s 35ms/step - loss: 0.2663 - acc: 0.8979 - val_loss: 0.2989 - val_acc: 0.8816
Epoch 3/20
30/30 [=====] - 1s 31ms/step - loss: 0.1894 - acc: 0.9293 - val_loss: 0.2992 - val_acc: 0.8801
Epoch 4/20
30/30 [=====] - 1s 36ms/step - loss: 0.1461 - acc: 0.9473 - val_loss: 0.3394 - val_acc: 0.8705
Epoch 5/20
30/30 [=====] - 1s 42ms/step - loss: 0.1121 - acc: 0.9625 - val_loss: 0.3741 - val_acc: 0.8688
Epoch 6/20
30/30 [=====] - 1s 41ms/step - loss: 0.0868 - acc: 0.9732 - val_loss: 0.4216 - val_acc: 0.8659
Epoch 7/20
30/30 [=====] - 1s 30ms/step - loss: 0.0650 - acc: 0.9803 - val_loss: 0.4729 - val_acc: 0.8668
Epoch 8/20
30/30 [=====] - 1s 30ms/step - loss: 0.0442 - acc: 0.9886 - val_loss: 0.5303 - val_acc: 0.8623
Epoch 9/20
30/30 [=====] - 1s 31ms/step - loss: 0.0273 - acc: 0.9958 - val_loss: 0.5880 - val_acc: 0.8595
Epoch 10/20
30/30 [=====] - 1s 34ms/step - loss: 0.0167 - acc: 0.9985 - val_loss: 0.6399 - val_acc: 0.8584
Epoch 11/20
30/30 [=====] - 1s 35ms/step - loss: 0.0102 - acc: 0.9996 - val_loss: 0.6888 - val_acc: 0.8585
Epoch 12/20
30/30 [=====] - 1s 30ms/step - loss: 0.0065 - acc: 0.9998 - val_loss: 0.7375 - val_acc: 0.8563
Epoch 13/20
30/30 [=====] - 1s 30ms/step - loss: 0.0045 - acc: 0.9999 - val_loss: 0.7758 - val_acc: 0.8560
Epoch 14/20
30/30 [=====] - 1s 30ms/step - loss: 0.0033 - acc: 0.9999 - val_loss: 0.8102 - val_acc: 0.8552
Epoch 15/20
30/30 [=====] - 1s 26ms/step - loss: 0.0026 - acc: 0.9999 - val_loss: 0.8360 - val_acc: 0.8569
Epoch 16/20
30/30 [=====] - 1s 25ms/step - loss: 0.0020 - acc: 0.9999 - val_loss: 0.8679 - val_acc: 0.8555
Epoch 17/20
30/30 [=====] - 1s 30ms/step - loss: 0.0015 - acc: 0.9999 - val_loss: 0.8976 - val_acc: 0.8567
Epoch 18/20
30/30 [=====] - 1s 48ms/step - loss: 0.0012 - acc: 1.0000 - val_loss: 0.9259 - val_acc: 0.8563
```

```

Epoch 19/20
30/30 [=====] - 1s 43ms/step - loss: 9.5346e-04 - acc: 1.0000 - val_loss: 0.9568 - val_acc: 0.8561
Epoch 20/20
30/30 [=====] - 1s 31ms/step - loss: 7.5277e-04 - acc: 1.0000 - val_loss: 0.9860 - val_acc: 0.8558
The Dense Convolutional Neural Network 2 layers took 21.9591 seconds to train.

```

```

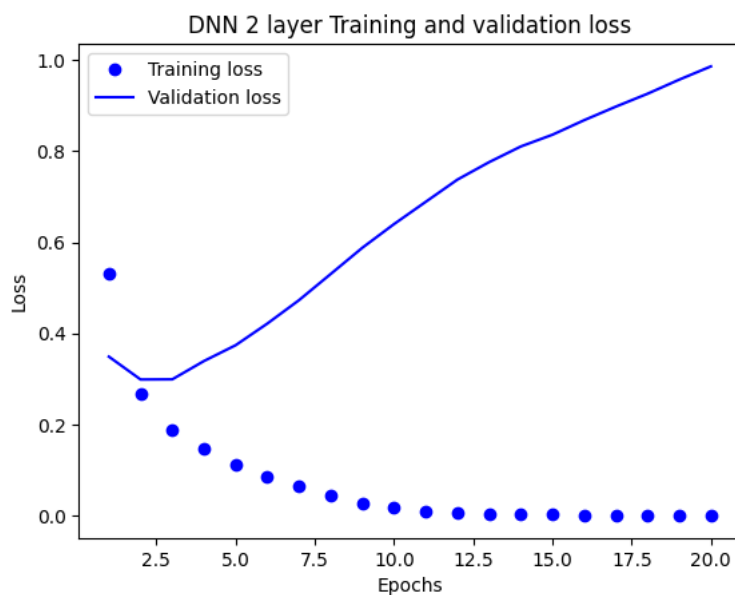
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

# "bo" is for "blue dot"
plt.plot(epochs, loss, 'bo', label='Training loss')
# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('DNN 2 layer Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```




```

plt.clf() # clear figure
acc_values = history_dict['acc']
val_acc_values = history_dict['val_acc']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('DNN 2 layer Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```



```
model2.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
dense_3 (Dense)	(None, 32)	160032
dense_4 (Dense)	(None, 32)	1056
dense_5 (Dense)	(None, 32)	1056
dense_6 (Dense)	(None, 1)	33
=====		
Total params: 162,177		
Trainable params: 162,177		
Non-trainable params: 0		

```
from numpy.ma.core import argmax
pred = model2.predict(x_test)
classes_x=argmax(pred,axis=-1)
#accuracy
accuracy_score(y_test,classes_x)
```

```
782/782 [=====] - 2s 2ms/step
0.5
```